

Python Tutorial on Implementation of Gradient Descent Optimization algorithms from Scratch

Example : Implementation of Gradient Descent Algorithms from scratch in Python

Table of Contents

1. [Create Data](#)
2. [Solve Directly](#)
3. [Batch Gradient Descent](#)
4. [Stochastic Gradient Descent](#)
5. [Mini-batch Gradient Descent](#)

In [1]:



```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 %matplotlib inline
```

1.Create Data

```
In [15]: ▶ 1 # generate random data with some noise
2
3 X = 2 * np.random.rand(100, 1)
4 print(X.shape)
5 # print(X)
6 y = 8 +(3*X) + (np.random.randn(100, 1))
7 print(y.shape)
8
9 ## For printing together input and target
10 X_y=np.c_[X,y]
11 print("Input and target")
12 print(X_y)
```

(100, 1)

(100, 1)

Input and target

```
[[ 0.04497682  9.10423587]
 [ 0.44369684  9.91643706]
 [ 1.14757428 11.57207686]
 [ 1.98332104 12.02308413]
 [ 0.03038902  9.2510614 ]
 [ 1.06406009 10.54856201]
 [ 0.76146702 11.78215315]
 [ 0.92371199  9.46087616]
 [ 1.14317242 11.56766946]
 [ 1.55185131 12.12205791]
 [ 0.0186962   7.2777673 ]
 [ 1.1789661   11.50784111]
 [ 0.12408403  7.41553359]
 [ 1.19483023 12.57702559]
 [ 1.59064174 14.23886903]
 [ 1.40627822 12.4646239 ]
 [ 0.41236386  9.66891799]
 [ 0.01554289  7.50275661]
 [ 1.24460193 12.63655415]
 [ 0.59811763 11.0568907 ]
 [ 0.41969231  9.65690536]
 [ 1.36654913 12.13159462]
 [ 1.64921642 13.65952935]
 [ 0.35949318  6.70617922]
 [ 0.02194567  8.71344024]
 [ 0.19417357  7.8774343 ]
 [ 0.43027749  8.26604118]
 [ 0.01570807  7.77699381]
 [ 1.8942345   14.27457547]
 [ 1.84577681 13.78644959]
 [ 1.0891461   11.9357483 ]
 [ 0.94948277 10.53574713]
 [ 1.8847801   13.73178354]
 [ 1.26666735 11.90012728]
 [ 0.40107006  7.84937596]
 [ 1.44085866 11.71417226]
 [ 1.61521082 13.30733268]
 [ 1.50202438 11.5954719 ]]
```

[0.88315884 11.46431646]
[0.47911189 9.47640518]
[0.30588616 10.13541128]
[1.52129161 12.84197568]
[0.25075942 7.88154753]
[0.90823295 10.79025615]
[1.63317685 12.44083037]
[1.68238394 14.15681462]
[1.36562465 12.50388206]
[0.77130311 12.89670151]
[1.10717202 10.36614913]
[0.23961462 7.03771564]
[0.55449143 8.97155405]
[1.96119551 13.11300052]
[1.02226458 10.56326872]
[0.13509331 9.68292579]
[0.27385538 7.45643523]
[0.68252181 10.8104117]
[0.47555176 9.06428252]
[1.09838153 11.87228133]
[1.08531632 11.92645865]
[0.75247426 8.81934867]
[1.99022675 15.21941186]
[1.99980551 12.91520817]
[1.5864137 12.94750455]
[1.66776477 13.05471928]
[0.39009362 9.5872106]
[0.06842432 6.52626061]
[0.49160231 8.96437506]
[1.89122417 13.97621451]
[0.29577358 9.39384895]
[1.21440789 11.27633599]
[1.79309467 11.8779532]
[1.63039219 12.65629986]
[0.81381444 11.28511091]
[1.0785014 11.64870513]
[0.65317283 10.92945199]
[1.39720473 10.49816875]
[0.72273399 10.0385813]
[1.56516876 11.22598596]
[1.79435949 13.69041692]
[0.92515284 10.69258089]


```
[ 0.92666572 11.16743518]
[ 1.79482058 14.3119088 ]
[ 0.83466322 10.25302971]
[ 0.89994989  9.90721344]
[ 0.05247682  7.3231602 ]
[ 1.34408815 12.47571037]
[ 1.32256478 11.96633382]
[ 0.9896771  11.92459708]
[ 0.38734407  9.25595385]
[ 0.62419774  9.34182876]
[ 1.03598244  9.13251898]
[ 1.52120598 12.7770783 ]
[ 0.16201475  7.62316748]
[ 1.5294854  12.93170731]
[ 0.77389001 11.13815733]
[ 1.43854668 14.43331059]
[ 0.77970708  9.08175759]
[ 0.42117791  9.99767758]
[ 0.38189651  8.82041598]
[ 0.45969414 11.29653945]]
```

In [16]: ▶

```
1 # X=2*np.random.rand(100,1)
2 # print(X)
3 # print(X.shape)
4 # y=8+3*X+np.random.rand(100,1)
5 # print(y.shape)
```

```
In [17]: ▶ 1 # Scatter plot to see the relation between X and y
          2 plt.figure(figsize=(15,10))
          3 plt.scatter(X, y, marker='x')
          4 plt.xlabel('X')
          5 plt.ylabel('y', rotation=0)
```

```
Out[17]: Text(0, 0.5, 'y')
```



2. Solve using Numpy Library

$$\nabla_w \hat{L}(f_w) = \nabla_w \frac{1}{n} \|Xw - y\|_2^2 = 0$$

```
In [18]: ▶ 1 # # Understanding numpy functions
2 # import numpy as np
3 # X = np.arange(6)
4 # print(X)
5 # X = X.reshape((6, 1))
6 # print(X)
7 # X1=np.ones_like(X)           # Return an array of ones with the same shape and type as a given array.
8 # print(X1)
9 # z=np.c_[X1,X]               # Translates slice objects to concatenation along the second axis.
10 # print(z)
```

```
In [19]: ▶ 1 # Adding bias unit to every vector in X
2 B=np.ones_like(X)
3 X_b=np.c_[B,X]
4 print(X_b)
5 print(X_b.shape)
6
7 #y=y.reshape((6,2))
8 theta=np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y)
9 print(theta)
```

[1.	0.04497682]
[1.	0.44369684]
[1.	1.14757428]
[1.	1.98332104]
[1.	0.03038902]
[1.	1.06406009]
[1.	0.76146702]
[1.	0.92371199]
[1.	1.14317242]
[1.	1.55185131]
[1.	0.0186962]
[1.	1.1789661]
[1.	0.12408403]
[1.	1.19483023]
[1.	1.59064174]
[1.	1.40627822]
[1.	0.41236386]
[1.	0.01554289]
[1.	1.24460193]
[1.	0.59811763]
[1.	0.41969231]
[1.	1.36654913]
[1.	1.64921642]
[1.	0.35949318]
[1.	0.02194567]
[1.	0.19417357]
[1.	0.43027749]
[1.	0.01570807]
[1.	1.8942345]
[1.	1.84577681]
[1.	1.0891461]
[1.	0.94948277]
[1.	1.8847801]
[1.	1.26666735]
[1.	0.40107006]
[1.	1.44085866]
[1.	1.61521082]
[1.	1.50202438]
[1.	0.88315884]
[1.	0.47911189]
[1.	0.30588616]

[1.	1.52129161]
[1.	0.25075942]
[1.	0.90823295]
[1.	1.63317685]
[1.	1.68238394]
[1.	1.36562465]
[1.	0.77130311]
[1.	1.10717202]
[1.	0.23961462]
[1.	0.55449143]
[1.	1.96119551]
[1.	1.02226458]
[1.	0.13509331]
[1.	0.27385538]
[1.	0.68252181]
[1.	0.47555176]
[1.	1.09838153]
[1.	1.08531632]
[1.	0.75247426]
[1.	1.99022675]
[1.	1.99980551]
[1.	1.5864137]
[1.	1.66776477]
[1.	0.39009362]
[1.	0.06842432]
[1.	0.49160231]
[1.	1.89122417]
[1.	0.29577358]
[1.	1.21440789]
[1.	1.79309467]
[1.	1.63039219]
[1.	0.81381444]
[1.	1.0785014]
[1.	0.65317283]
[1.	1.39720473]
[1.	0.72273399]
[1.	1.56516876]
[1.	1.79435949]
[1.	0.92515284]
[1.	0.92666572]
[1.	1.79482058]
[1.	0.83466322]

```
[1.      0.89994989]
[1.      0.05247682]
[1.      1.34408815]
[1.      1.32256478]
[1.      0.9896771 ]
[1.      0.38734407]
[1.      0.62419774]
[1.      1.03598244]
[1.      1.52120598]
[1.      0.16201475]
[1.      1.5294854 ]
[1.      0.77389001]
[1.      1.43854668]
[1.      0.77970708]
[1.      0.42117791]
[1.      0.38189651]
[1.      0.45969414]]
(100, 2)
[[7.87752609]
 [3.11459601]]
```

In [20]: ▶

```
1 # Adding bias unit to every vector in X
2 B= np.ones_like(X)           # Return an array of ones with the same shape and type as a given array.
3 X_b = np.c_[B, X]           # Translates slice objects to concatenation along the second axis.
4 # X_b = X
5 print(X_b)
6 theta = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y) # Compute the (multiplicative) inverse of a matrix.
7
8 # Values are very close to are thetas 8 and 3
9 print(theta)
10
```

[1.	0.04497682]
[1.	0.44369684]
[1.	1.14757428]
[1.	1.98332104]
[1.	0.03038902]
[1.	1.06406009]
[1.	0.76146702]
[1.	0.92371199]
[1.	1.14317242]
[1.	1.55185131]
[1.	0.0186962]
[1.	1.1789661]
[1.	0.12408403]
[1.	1.19483023]
[1.	1.59064174]
[1.	1.40627822]
[1.	0.41236386]
[1.	0.01554289]
[1.	1.24460193]
[1.	0.59811763]
[1.	0.41969231]
[1.	1.36654913]
[1.	1.64921642]
[1.	0.35949318]
[1.	0.02194567]
[1.	0.19417357]
[1.	0.43027749]
[1.	0.01570807]
[1.	1.8942345]
[1.	1.84577681]
[1.	1.0891461]
[1.	0.94948277]
[1.	1.8847801]
[1.	1.26666735]
[1.	0.40107006]
[1.	1.44085866]
[1.	1.61521082]
[1.	1.50202438]
[1.	0.88315884]
[1.	0.47911189]
[1.	0.30588616]

[1.	1.52129161]
[1.	0.25075942]
[1.	0.90823295]
[1.	1.63317685]
[1.	1.68238394]
[1.	1.36562465]
[1.	0.77130311]
[1.	1.10717202]
[1.	0.23961462]
[1.	0.55449143]
[1.	1.96119551]
[1.	1.02226458]
[1.	0.13509331]
[1.	0.27385538]
[1.	0.68252181]
[1.	0.47555176]
[1.	1.09838153]
[1.	1.08531632]
[1.	0.75247426]
[1.	1.99022675]
[1.	1.99980551]
[1.	1.5864137]
[1.	1.66776477]
[1.	0.39009362]
[1.	0.06842432]
[1.	0.49160231]
[1.	1.89122417]
[1.	0.29577358]
[1.	1.21440789]
[1.	1.79309467]
[1.	1.63039219]
[1.	0.81381444]
[1.	1.0785014]
[1.	0.65317283]
[1.	1.39720473]
[1.	0.72273399]
[1.	1.56516876]
[1.	1.79435949]
[1.	0.92515284]
[1.	0.92666572]
[1.	1.79482058]
[1.	0.83466322]


```

[1.      0.89994989]
[1.      0.05247682]
[1.      1.34408815]
[1.      1.32256478]
[1.      0.9896771 ]
[1.      0.38734407]
[1.      0.62419774]
[1.      1.03598244]
[1.      1.52120598]
[1.      0.16201475]
[1.      1.5294854 ]
[1.      0.77389001]
[1.      1.43854668]
[1.      0.77970708]
[1.      0.42117791]
[1.      0.38189651]
[1.      0.45969414]]
[[7.87752609]
 [3.11459601]]

```

3. Batch Gradient Decent Algorithm

Below mentioned **Cost Function** which stands for **Mean Squared Error(MSE)** is for **Linear Regression** and **Gradient** is derived from the it.

Cost

$$J(\theta) = 1/m \sum_{i=1}^m (h(\theta)^{(i)} - y^{(i)})^2$$

Gradient

$$\frac{\partial J(\theta)}{\partial \theta_j} = 2/m \sum_{i=1}^m (h(\theta)^{(i)} - y^{(i)}) \cdot X_j^{(i)}$$

Below mentioned **Cost Function** which stands for **Mean Squared Error(MSE)** is for **Linear Regression** and **Gradient** is derived from the it.

Cost

$$J(\theta) = 1/2m \sum_{i=1}^m (h(\theta)^{(i)} - y^{(i)})^2$$

Gradient

$$\frac{\partial J(\theta)}{\partial \theta_j} = 1/m \sum_{i=1}^m (h(\theta)^{(i)} - y^{(i)}) \cdot X_j^{(i)}$$

Function for computing Cost function

```
In [21]: ► 1 def cal_cost(theta, X, y):
2         '''
3         Calculates the cost for given X and y.
4
5         Parameters
6         -----
7         theta: Vector of thetas
8         X: Row of X's
9         y: Actual y's
10
11         Returns
12         -----
13         Calculated cost
14         '''
15
16         m = len(y)
17         predictions = np.dot(X, theta)
18         cost = (1/ m) * np.sum(np.square(predictions - y))
19
20         return cost
```

Function for Batch gradient Decent Algorithm

```
In [22]: ▶ 1 def gradient_descent(X, y, theta, learning_rate, iterations):
2     '''
3     Parameters
4     -----
5     X: Matrix of X with bias units added to every vector of X
6     y: Vector of y
7     theta: Vector initialized randomly
8     learning_rate: learning rate that will be used as step size
9     iterations = Number of iterations
10
11     Returns
12     -----
13     The final theta vector, array of cost and theta history over no of iterations
14     '''
15
16     m = len(y)
17     cost_history = np.zeros(iterations)
18     theta_history = np.zeros((iterations, 2))
19
20     for i in range(iterations):
21         #Forward propagation
22         predictions = np.dot(X, theta)
23         #Compute Cost (Value of Loss function)
24         cost_history[i] = cal_cost(theta, X, y)
25         # Backpropogation
26         grad =(2/m) * (X.T.dot((predictions - y)))
27         # Update Parameters
28         theta = theta - (learning_rate * grad)
29         theta_history[i, :] = theta.T
30
31
32     return theta, cost_history, theta_history
```

Calling function

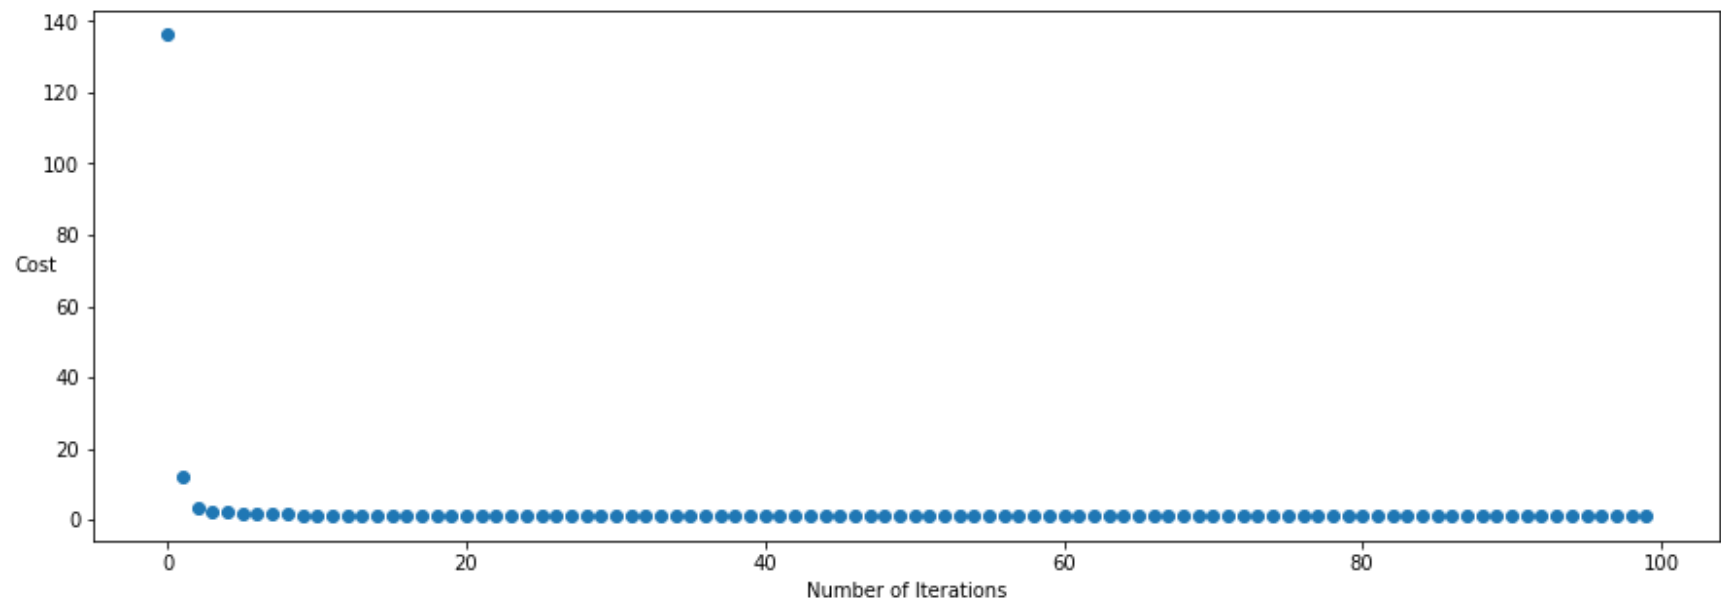
```
In [23]: ▶ 1 # Using gradient descent to find out the thetas between our X and y relation
2 import time
3 lr = 0.3
4 n_iter = 100
5
6 theta = np.random.randn(2,1)
7 print("Initial value of weights")
8 print(theta)
9 X_b = np.c_[np.ones((len(X),1)),X]
10
11 ## Calling BGD function
12 t10 = time.time()
13 theta,cost_history,theta_history = gradient_descent(X_b, y, theta, lr, n_iter)
14 t11 = time.time()
15 time_BGD = t11 - t10
16 print("Processing Time of MBG Algorithm")
17 print(time_BGD)
18
19
20 # Predicted value by NNs
21 y_predicted= np.dot(X_b, theta)
22 print("Optimal value of weights")
23 print('{:<10}{:.3f}'.format('Theta0:',theta[0][0]))
24 print('{:<10}{:.3f}'.format('Theta1:',theta[1][0]))
25 print("Minimum value of cost function")
26 print('{:<10}{:.3f}'.format('Cost/MSE:',cost_history[-1]))
27
28 # print(y_predicted)
29
30 predicted__actual = np.c_[y_predicted, y]
31 print("Comparing predicted vs actual")
32 print(predicted__actual)
```

[5.2555500 5.51045700]
[11.45175653 11.57207686]
[14.0548556 12.02308413]
[7.97206146 9.2510614]
[11.19163502 10.54856201]
[10.24914889 11.78215315]
[10.75449303 9.46087616]
[11.43804608 11.56766946]
[12.71095756 12.12205791]
[7.93564186 7.2777673]
[11.5495326 11.50784111]
[8.26389315 7.41553359]
[11.59894458 12.57702559]
[12.83177803 14.23886903]
[12.25754129 12.4646239]
[9.16179784 9.66891799]
[7.92582023 7.50275661]
[11.75396842 12.63655415]
[9.74036482 11.0568907]
[9.18462378 9.65690536]
[10.40270715 10.40150460]




```
In [24]: 1 # plotting the cost history vs number of iterations
2
3 plt.figure(figsize=(15,5))
4 plt.scatter(range(n_iter), cost_history)
5 plt.xlabel('Number of Iterations')
6 plt.ylabel('Cost', rotation=0)
7
```

Out[24]: Text(0, 0.5, 'Cost')



```
In [25]: 1 # plot prediction line
2 plt.figure(figsize=(15,5))
3 plt.scatter(X, y, marker='x')
4 plt.plot(X, y_predicted, 'red')
5 plt.xlabel('X')
6 plt.ylabel('y', rotation=0)
```

Out[25]: Text(0, 0.5, 'y')



4. Stochastic Gradient Descent (SGD) Algorithm


```

In [26]: 1 def stochastic_gradient_descent(X, y, theta, learning_rate, iterations):
2         '''
3         Parameters
4         -----
5         X: Matrix of X with bais units added to every vector of X
6         y: Vector of y
7         theta: Vector initialized randomly
8         learning_rate: learning rate that will be used as step size
9         iterations = Number of iterations
10
11         Returns
12         -----
13         The final theta vector, array of cost and theta history over no of iterations
14         '''
15         m = len(y)
16         theta_history = np.zeros((iterations, 2))
17         cost_history = np.zeros(iterations)
18         # theta_history = []
19
20         for i in range(iterations):
21             cost_per_iteration = .0
22             for j in range(m):
23                 # Step : Shuffle (X,y)
24                 X_rand_idx = np.random.randint(m)
25                 X_inner = X[X_rand_idx].reshape(1, X.shape[1])
26                 y_inner = y[X_rand_idx].reshape(1, 1)
27
28                 # Forward Propgation
29                 predictions = np.dot(X_inner, theta)
30
31                 # Computing cost value
32                 cost_per_iteration += cal_cost(theta, X_inner, y_inner)
33
34                 # Backpropagation
35                 grad = (2/m) * (X_inner.T.dot((predictions - y_inner)))
36
37                 # Weight updation
38                 theta = theta - (learning_rate * grad)
39
40             cost_history[i] = cost_per_iteration
41

```

42

43

```
return theta, cost_history, theta_history
```

In [31]: ▶

```
1  # Using stochastic gradient descent to find out the thetas between our X and y relation
2  import time
3  lr = 0.3
4  n_iter = 100
5  theta = np.random.randn(2,1)
6  print("Initial value of weights")
7  print(theta)
8
9  X_b = np.c_[np.ones((len(X),1)),X]
10
11  ## calling SGD
12  t20 = time.time()
13  theta,cost_history,theta_history = stochastic_gradient_descent(X_b, y, theta, lr, n_iter)
14  t21 = time.time()
15  time_SGD = t21 - t20
16  print("Processing Time of SGD Algorithm")
17  print(time_SGD)
18  # theta,cost_history,theta_history = stochastic_gradient_descent(X_b, y, theta, lr, n_iter)
19  print(theta)
20
21
22  # Predicted value by NNs
23  y_predicted= np.dot(X_b, theta)
24
25  print("Optimal value of weights")
26  print('{:<10}{:.3f}'.format('Theta0:',theta[0][0]))
27  print('{:<10}{:.3f}'.format('Theta1:',theta[1][0]))
28  print("Minimum value of cost function")
29  print('{:<10}{:.3f}'.format('Cost/MSE:',cost_history[-1]))
30
31
32  # print(y_predicted)
33  predicted__actual = np.c_[y_predicted, y]
34  print("Comparing predicted vs actual")
35  print(predicted__actual)
36
37
```

```
[ 9.26269015  8.26604118]
[ 7.9709022   7.77699381]
[13.82434371 14.27457547]
[13.67335075 13.78644959]
[11.31570832 11.9357483 ]
[10.88052087 10.53574713]
[13.79488402 13.73178354]
[11.86886012 11.90012728]
[ 9.17168056  7.84937596]
[12.41163588 11.71417226]
[12.95491287 13.30733268]
[12.60222674 11.5954719 ]
[10.67385716 11.46431646]
[ 9.41485695  9.47640518]
[ 8.87508987 10.13541128]
[12.66226294 12.84197568]
[ 8.70331633  7.88154753]
[10.75198745 10.79025615]
[13.01089457 12.44083037]
[13.16422264 14.15681462]
```




In [32]: ►

```
1 # def stochastic_gradient_descent(X, y, theta, Learning_rate, iterations):
2 #     '''
3 #         Parameters
4 #         -----
5 #         X: Matrix of X with bias units added to every vector of X
6 #         y: Vector of y
7 #         theta: Vector initialized randomly
8 #         Learning_rate: Learning rate that will be used as step size
9 #         iterations = Number of iterations
10
11 #     Returns
12 #     -----
13 #     The final theta vector, array of cost and theta history over no of iterations
14 #     '''
15
16 #     m = len(y)
17 #     theta_history = np.zeros((iterations, 2))
18 #     cost_history = np.zeros(iterations)
19 #     theta_history = []
20
21 #     for i in range(iterations):
22 #         cost_per_iteration = .0
23 #         for j in range(m):
24 #             # Step : Shuffle (X,y)
25 #             X_rand_idx = np.random.randint(m)
26 #             X_inner = X[X_rand_idx].reshape(1, X.shape[1])
27 #             y_inner = y[X_rand_idx].reshape(1, 1)
28
29 #             predictions = np.dot(X_inner, theta)
30
31 #             theta = theta - (2/m) * Learning_rate * (X_inner.T.dot((predictions - y_inner)))
32 #             cost_per_iteration += cal_cost(theta, X_inner, y_inner)
33 #             cost_history[i] = cost_per_iteration
34
35 #     return theta, cost_history, theta_history
```


```
In [20]: 1 # plotting the cost history vs number of iterations
2
3 plt.figure(figsize=(15,5))
4 plt.scatter(range(n_iter), cost_history)
5 plt.xlabel('Number of Iterations')
6 plt.ylabel('Cost', rotation=0)
7
```

Out[20]: Text(0, 0.5, 'Cost')



```
In [20]: 1 # plot prediction line
2 plt.figure(figsize=(15,5))
3 plt.scatter(X, y, marker='x')
4 plt.plot(X, np.dot(X_b, theta), 'red')
5 plt.xlabel('X')
6 plt.ylabel('y', rotation=0)
```

Out[20]: Text(0, 0.5, 'y')



5. Mini-batch Gradient Descent (MBGD) Algorithm


```

In [27]: 1 def mini_batch_gradient_descent(X, y, theta, learning_rate, iterations):
2         '''
3         Parameters
4         -----
5         X: Matrix of X with bias units added to every vector of X
6         y: Vector of y
7         theta: Vector initialized randomly
8         learning_rate: learning rate that will be used as step size
9         iterations = Number of iterations
10
11         Returns
12         -----
13         The final theta vector, array of cost and theta history over no of iterations
14         '''
15
16         m = len(y)
17         cost_history = np.zeros(iterations)
18         # print(cost_history)
19         theta_history = np.zeros((iterations, 2))
20         # print(theta_history)
21         # Batch size
22         batch_size=10
23         for i in range(iterations):
24             cost_per_iteration = .0
25             # Step 1: Shuffle (X,y)
26             rand_indicies = np.random.permutation(m)
27             X = X[rand_indicies]
28             y = y[rand_indicies]
29
30             for j in range(0, m, batch_size):
31                 X_inner = X[i:i+batch_size]
32                 y_inner = y[i:i+batch_size]
33                 # forward propagation
34                 predictions = np.dot(X_inner, theta)
35                 # weight updation equation
36                 # theta = theta - (2/m) * learning_rate * (X_inner.T.dot((predictions - y_inner)))
37                 # cost_per_iteration += cal_cost(theta, X_inner, y_inner)
38                 # Computing cost value
39                 cost_per_iteration += cal_cost(theta, X_inner, y_inner)
40                 # Backpropagation
41                 grad= (2/m) * (X_inner.T.dot((predictions - y_inner)))

```

```
42         # Weight updation
43         theta = theta - (learning_rate * grad)
44         cost_history[i] = cost_per_iteration
45
46     return theta, cost_history, theta_history
```

In [28]: ▶


```
1 # Using mini batch gradient descent to find out the thetas between our X and y relation
2 lr = 0.3
3 n_iter = 100
4 theta = np.random.randn(2,1)
5 print("Initial value of weights")
6 print(theta)
7
8 # print(X_b)
9 X_b = np.c_[np.ones((len(X),1)),X]
10 # print(X_b)
11
12 ## calling MBGD
13 t30 = time.time()
14 theta,cost_history,theta_history = mini_batch_gradient_descent(X_b, y, theta, lr, n_iter)
15 t31 = time.time()
16 time_MBGD = t31 - t30
17 print("Processing Time of MBGD Algorithm")
18 print(time_MBGD)
19 # theta,cost_history,theta_history = mini_gradient_descent(X_b, y, theta, lr, n_iter)
20 # print(theta)
21
22 # Predicted value by NNs
23 y_predicted= np.dot(X_b, theta)
24
25 print("Optimal value of weights")
26 print('{:<10}{:.3f}'.format('Theta0:',theta[0][0]))
27 print('{:<10}{:.3f}'.format('Theta1:',theta[1][0]))
28 print("Minimum value of cost function")
29 print('{:<10}{:.3f}'.format('Cost/MSE:',cost_history[-1]))
30
31 # print(y_predicted)
32 predicted__actual = np.c_[y_predicted, y]
33 print("Comparing predicted vs actual")
34 print(predicted__actual)
```

[12.12343891 12.13159462]
[13.04414139 13.65952935]
[8.84326137 6.70617922]
[7.74380331 8.71344024]
[8.30478315 7.8774343]
[9.07381964 8.26604118]
[7.72348624 7.77699381]
[13.84221304 14.27457547]
[13.68437689 13.78644959]
[11.21988316 11.9357483]
[10.76497248 10.53574713]
[13.8114182 13.73178354]
[11.79810449 11.90012728]
[8.97868538 7.84937596]
[12.36547953 11.71417226]
[12.93337851 13.30733268]
[12.5647082 11.5954719]
[10.5489425 11.46431646]
[9.23288283 9.47640518]
[8.66865284 10.13541128]




```
In [29]: 1 # plotting the cost history vs number of iterations
2
3 plt.figure(figsize=(15,5))
4 plt.scatter(range(n_iter), cost_history)
5 plt.xlabel('Number of Iterations')
6 plt.ylabel('Cost', rotation=0)
7
```

Out[29]: Text(0, 0.5, 'Cost')



```
In [30]: 1 # plot prediction line
2
3 plt.figure(figsize=(15,5))
4 plt.scatter(X, y, marker='x')
5 plt.plot(X, np.dot(X_b, theta), 'red')
6 plt.xlabel('X')
7 plt.ylabel('y', rotation=0)
```

Out[30]: Text(0, 0.5, 'y')




Comparison of Processing Speed of above three Gradient Decent Optimization Algorithms Implementation

```

In [33]: ► 1 import matplotlib.pyplot as plt
2 %matplotlib inline
3 # plt.style.use('ggplot')
4
5 x = ['Method-1', 'Method-2', 'Method-3']
6 Processing_Times = [time_BGD, time_SGD, time_MBGD]
7
8
9 x_pos = [i for i, _ in enumerate(x)]
10
11 plt.bar(x_pos, Processing_Times, color='green')
12 plt.xlabel("<-----Method-----> ")
13 plt.ylabel("<-----Time-----> ")
14 plt.title(" Time Vs Method ")
15
16 plt.xticks(x_pos, x)
17
18 plt.show()

```



In []: ▶

1